




The Tableau Plateau

From dashboards to AI-ready data products


Data vs Vibes



Data vs Vibes



Data vs Vibes




Data vs Vibes

vibes are the most
important KPI 

Data vs Vibes

vibes are the most
important KPI



Data vs. Vibes



Humanities 101

Data vs. Vibes



The Suffolk Journey

Suffolk's Path to Enlightenment

Who is Suffolk?


Boston Based top-50 General Contractor (GC)

Wow - I don't think of construction as a tech-forward industry!

- Construction productivity in **2023** remained roughly where it was in **1948**.
- Only **12%** of industry respondents used digital tools across all projects.
- **1.4%** of construction firms used AI, versus **18.1%** in information industries.

Suffolk has been at the forefront of driving adoption


- Investing in ConTech for ~20 years
- First CDO in the industry



Suffolk's Path to Enlightenment



“Every business unit needs at least one dashboard.”
— John Fish (CEO), ~2018



Three teams, one deliverable — by chance


Data Science

Data Analysis

Data Viz


Which team you talked to decided your solution — not the problem.

Three teams, one deliverable — by chance



Which team you talked to decided your solution — not the problem.

A dashboard is a terminal shape



- You can make it prettier
- You can make it faster
- But its still always a dashboard
- You can't make it an alert, or an API, or a voice update or a judgement at the point of work.

Ask a better question

The metric question

Has this earned the right
to be codified?

The vehicle question

Is a dashboard even the
right shape for this?

THE METRIC QUESTION

Problem: dashboards arrive too early


“Every business unit needs at least one dashboard.”

A useful mandate became an unquestioned ritual

“Data driven” became shorthand for “build a dashboard”



Symptom 1: the kitchen sink



THE METRIC QUESTION

Symptom 1: the kitchen sink



THE METRIC QUESTION

Symptom 1: the kitchen sink



Symptom 2: sprawl

From: [REDACTED]

Sent: Thursday, July 2, 2020 11:17 AM

To: [REDACTED]

Subject: Dashboard Inventory

Hi [REDACTED]

As we head into business planning, I know many people across our teams have the mandate to streamline our dashboards and refine our metrics to better align with our highest-priority objectives.

As part of this effort, [REDACTED] would like to review an inventory of our current dashboards, including:

- The dashboards that are currently in use
- The metrics tracked within each dashboard
- Business owner
- Primary audience

If it's feasible, it would be great to have this pulled together for him early next week. If that timing is challenging, just let me know what you think is realistic.

THE METRIC QUESTION

Symptom 2: sprawl

From: [REDACTED]

Sent: Thursday, July 2, 2026 11:17 AM

To: [REDACTED]

Subject: Dashboard Inventory

Hi [REDACTED]

As we head into business planning, I know many people across our teams have the mandate to streamline our dashboards and refine our metrics to better align with our highest-priority objectives.

As part of this effort, [REDACTED] would like to review an inventory of our current dashboards, including:



- The dashboards that are currently in use
- The metrics tracked within each dashboard
- Business owner
- Primary audience

If it's feasible, it would be great to have this pulled together for him early next week. If that timing is challenging, just let me know what you think is realistic.



THE METRIC QUESTION

Symptom 2: sprawl



One team, many deliverables



Blended team with Product Management support.

Structured Guided Intake

Oca Finance wants a dashboard combining all subcontractor compliance information.

Idea Evaluation

Problem

What is the problem, and why does it matter? What happens if we do nothing?

Opportunity & Value

How could we improve this? What value or capability is unlocked for the business?

Desirability (Business Demand)

How strong is business demand? Who is asking for this and why now?

Feasibility / Readiness

Test	Response
Data Readiness Does system(s) exist, reliable data exist?	
Process Readiness Do there exist viable process flow / workflow?	
User Readiness Does the business accept / need?	
Technical Readiness Can we build this / build / build it?	
AI Readiness Can AI (GenAI, etc.) be leveraged?	
Scalability Potential Can this solve other use cases?	

Initiative Brief

Strategic Context

What is happening in the business today, and why does this matter now?

Decision & Capability Gap

Decision to Improve

What decision or workflow are we trying to improve?

Current Gap

- What is not working today?
- Where are the breakdowns (data, process, consistency, timing)?

Target Capability (Long-Term Vision)

What capability should exist if this is fully realized?

- What new behavior or workflow does this enable?
- How should the business operate differently?

Why Now (Priority)

Why should we invest in this now vs later?

- What makes this urgent or important today?
- What happens if we delay?

1. Overview

FieldValue

Product	
Technical Product Manager	
Data Strategy & Ops	
Executive Sponsor	
Team	

Description

[[Short description of the product, what it does, who it serves, and the core value proposition.]]

2. Problem Statement

[[Describe the core problem this product is solving.]]

This leads to: - (Pain point 1),

- (Pain point 2H),
- (Pain point 2H)

3. Goals & Non-Goals

Goals

- (Goal 1),
- (Goal 2),
- (Goal 3),

Non-Goals

- (Explicitly out-of-scope item 1),
- (Explicitly out-of-scope item 2H),
- (Explicitly out-of-scope item 3H)

Structured Guided Intake

Ops Finance wants a dashboard combining all subcontractor compliance information.

Idea Evaluation

Problem

What is the problem and why does it matter? What happens if we do nothing?

Opportunity & Value

How could we improve this? What value or capability is unlocked for the business?

Desirability (Business Demand)

How strong is business demand? Who is asking for this and why now?

Feasibility / Readiness

Test	Response
Data Readiness <i>Does systemized, reliable data exist?</i>	
Process Readiness <i>Is there a stable process this fits into?</i>	
User Readiness <i>Can the business adopt this?</i>	
Technical Readiness	

Structured Guided Intake

1. Overview

Field	Value
Product	
Technical Product Manager	
Data Strategy & Ops	
Executive Sponsor	
Team	

Description
{{Brief description of the product. What it does, who it serves, and the core value proposition.}}

2. Problem Statement

{{Describe the core problem this product is solving.}}

This leads to: - {{Pain point 1}}

- {{Pain point 2}}
- {{Pain point 3}}


3. Goals & Non-Goals

Goals

THE VEHICLE QUESTION

Problem: self-service is a myth

Even if the timing is right —
It still has to clear one more bar.



Most stakeholders don't want to hunt for insight —
they want it provided to them.

THE VEHICLE QUESTION

Why we can finally ask it: the ceiling is gone

Positron Assistant

Inside the IDE



Devin AI


Integrated into the workflow



The old excuse for defaulting to dashboards was speed. That excuse is gone.



Guardrails



**A new process +
Agentic Framework +
Custom Packages =
Unlock**

Our new bottleneck was now the ‘agile’ process



Guardrails

Spec Driven Development

How the Workflow Works

Each document type has a **template** (the structure) and a **skill** (the agent-assisted process for populating it). The skills are used via Devin. The template is never handed to an agent/blank. The skill runs an interview or reads source material and fills it in, tags every gap, and sets the status. A human then reviews and approves before the next step begins.



guided-discovery	The default first step for starting new work, unless a spec exists or a spec is explicitly requested	Bounded problem exploration (brief when the problem is already clear) that produces a discovery artifact feeding feature-spec creation via the spec skill; the SDD entry point
------------------	--	--

guided-discovery — the default first move on any new work, from the SuffolkPowers catalog.

A pattern library, without the pattern lock-in.

SuffolkPowers

obra/superpowers



Component	Location	Purpose
Devin Skills	<code>.devin/skills/</code>	Agent skills for common workflows (SDD, Databricks access, skill authoring)
Engineering Standards	<code>standards/</code>	Coding best practices, agent behavior guidelines, PR review expectations
SDD Playbooks	<code>playbooks/sdd/</code>	Spec-driven development workflow, templates, and execution guides
Skill Eval Suite	<code>tests/</code>	Automated testing framework for validating skill effectiveness
Setup Script	<code>setup.sh</code>	Bootstraps local Devin config from <code>suffolk-powers</code>

These skills are available in every Suffolk County public after-school program.

Skill	When to invoke	What it Does
using-sphinx-power	Start of any session in a Sphinx repo	Loads the imposed discipline and catalogs available skills
spec	Creating initial reports, communications plans, architecture specs	Doubles SDO workflow with templates and iteration patterns
characterize-read	Reading data from Describes in R or Python	Provides patterns for <code>describe()</code> (R) and <code>unlist()</code> (Python) usage
improve-sphinx-skills	Creating or adding skills in <code>sphinx-power</code> or <code>interactive-reports</code>	Introduces the skill structure, names, and source locations
eval-skill	Testing skill changes in a real or eval suite	Drill exercises automated skill evaluation and benchmarking
suffice-kazans	Writing or modifying R/Python Kazans (linea or fourfold theorems)	Defines what to log, at what level, and what to avoid (locanote / <code>print</code>)
lim-skill	Using LLM calls from <code>suffice</code> or <code>synapse-note</code>	Routes to approval wrappers (<code>execute::llm::chat()</code> / <code>python::llm::chat()</code>) and handles naming, logging, running, and structured output
lim-evaluation	Evaluating or regression-testing LLM outputs in an SDO workflow	Builds eval cases from a spec and routes to the approved grading engine
suffice-post-app-development	Developing or reviewing apps in the <code>suffice</code> (<code>spec</code> / <code>subject</code>) monorepo	Covers monorepo rules, language standards, libraries, deployment, and testing
test-rules	Creating, updating, or running tests for rule or agent producers	Covers when structure required, coverage strategies, sufficiency judgement, and the standard report for <code>ITF</code> Validities
feature-spec-execution-standards	Timing an approved feature spec and TD plan has a lean, test-first execution plan, after spec-review and before implementation	Maps every Mart-Kane acceptance criterion to a specific prescription and an implementation step, produces execution order and file ID approval
do-avoid-sub-debugging	Any bug, test failure, or unexpected behavior, before proposing a fix	Enforces root-cause-first debugging via four phases: strip test-case-tracing, remove-a-convrt, and <code>toolchain-based-waiting</code> references
prompt-to-narrow	The default first step for starting new work, unless a spec event or a spec is explicitly requested	Bounded system exploration (brief when the problem is already clear) that produces a discovery artifact feeding feature-spec creation via the spec skill, the SDO entry point
spec-quality-review	Indocentent read-only review of a feature spec, before execution planning	Flags ambiguity, missing acceptance criteria, scope gaps, and off-scriptal (misalignment + drift), produces spec-review and elegantly fits feature spec, later progresses to feature-spec-validation, planning

SuffolkPowers

obra/superpowers



Q

Component	Location	Purpose
Devin Skills	<code>.devin/skills/</code>	Agent skills for common workflows (SDD, Databricks access, skill authoring)
Engineering Standards	<code>standards/</code>	Coding best practices, agent behavior guidelines, PR review expectations
SDD Playbooks	<code>playbooks/sdd/</code>	Spec-driven development workflow, templates, and execution guides
Skill Eval Suite	<code>tests/</code>	Automated testing framework for validating skill effectiveness
Setup Script	<code>setup.sh</code>	Bootstraps local Devin config from <code>suffolk-powers</code>

Skill	When to Invoke	What It Does
using-satellite-powers	Start of any session in a satellite repo	Loads the execution discipline and catalogs available skills
tools	Creating a file/directory, terminalists home, structure specs	Guides SCD workflow with terminalists and extension patterns
statistics-read	Reading data from Descriptors in R or Python	Provides patterns for <code>summary()</code> and <code>info()</code> (Python) usage
artisan-satellite-skills	Creating or editing skills in satellite-powers or instance specs	Introduces skill structure, names, and scope decisions
eval-skill	Testing skill changes in a new eval sales	Defines <code>eval</code> automated skill evaluation and benchmarking
satellite-learn	Writing or modifying R/Python learning items in <code>learned</code> resources	Defines what to log, at what level, and what to avoid (background / python)
lrm-calls	Using LLM calls from Satellite R as <code>python</code> node	Routes to approved wrappers (<code>execute::llm::call / python::llm::call</code>) and marks programming, logging, naming, and structure subset
lrm-evaluation	Evaluating or regression-testing LLM outputs in an SCD workflow	Builds eval cases from a spec and routes to the approved grading engine
satellite-code-app-development	Developing or reviewing apps in the satellite-specs/ <code>business</code> monorepo	Covers recommended rules, approach standards, libraries, copy-paste, and testing
test-rules	Creating, updating, or running test for node in agent products	Covers when/where required, coverage strategies, efficiency judgement, and the standard report for TC Validities
feature-spec-execution-clarify	Turning an approved feature spec and TD plan into a lean, test-first execution plan, after some results review and before implementation	Vape every Mart-Kline acceptance criterion to as specific pre-creation and an implementation step, produces execution rules and for IC example
straw-man-debugging	Any bug, test failure, or unexpected behavior, before proposing a fix	Enforces root-cause-fix debugging via four phases: topic root-cause-finding, <code>define-as-clear</code> , and concludes based on calling references
product-to-spec	The DEBUT first spec for SAFING new work, unless a spec exists or a spec is explicitly requested	Bundles product evaluation brief when the problem is already clear that produces a discovery artifact feeding feature-spec creation via the spec skill, the SCD entry point
spec quality review	Independent read-only review of a feature spec before execution planning	Flags ambiguity, missing acceptance criteria, scope gaps, and off-script risk (misalignment + drift), produces <code>spec-review</code> and aligns the feature spec, faster progression to feature-spec-execution planning
incremental-implementation	Executing an IC-approved execution plan in an existing code - the last step of the SCD cycle, after feature-spec-execution-planning	The SCD execution engine consumes execution plan and runs as step at a time - using <code>test + implement + verify + self review + summary</code> ending in a PR. Enforces the re-implementation <code>without-as-approved-plan</code> gate
context engineering	Starting a new coding session, switching specs or a workflow setting up a new project, or when output quality degrades	Loads and initializes the five-layer working context (rules files, specs, source, modules for look, conversion history) in agent mode before acting in-session discipline, no artifact
branching-and-styling	Routing or styling any visual or interactive parts, tables, dashboards, Shiny apps, Quarto docs, presentations, CSS in R as <code>python</code>	Uses agents at the most convenient and the <code>isocritical / python</code> functions that apply it, instead of hardcoded colors, fonts, or logos

Getting off the Plateau



Different problems deserve different shapes

Stable, recurring monitoring	Dashboard
Exceptions requiring attention	Alert
Questions requiring exploration	Conversational

Weekly Safety Summary

ABOUT THIS REPORT

This is an automated email for the Safety Model. Incidents and Observations below describe each project this week vs its own recent weeks. Risk Factors compare each project against typically observed values across all project-weeks, showing the primary reasons contributing to a project's increased risk.

Tier Definitions

Tier 1: Significant impact on model's predictions
Tier 2: Moderate impact on model's predictions
Tier 3: Minor impact on model's predictions
Tier 4: Negligible impact on model's predictions

Non-Influencable Factors

Project metadata that cannot be changed: % complete, region, SBU, and interaction terms.

Further Investigation

[Safety Model Dashboard - Webpage Version](#)

999999 — Project XXXXX

Incidents Summary

2 in last 4 weeks · 1 recordable

Subcontractor YYY had two eyeface injuries within six days, both involving overhead electrical work with conduit and wire splicing. The recordable incident required medical treatment for a laceration above the right eye from overhead conduit contact.

Observations Summary

17 obs this week · 178.18 hours/cbs

- Risk profile elevated vs typical projects though stable here, with findings concentrated on fall protection, demolition work, and fire prevention across multiple trades.
- Highest-risk: worker not wearing fall protection in restricted zone on terrace.
- Fall protection – personal trending higher: workers in scissor lifts without PFAS, not tied off in restricted zones.
- Demolition safety, exposed surfaces / protrusions, perimeter protection / overhead hazard, and fire prevention / hot work safety also trending higher: coiling demo before electrical make-safe, uncovered slab penetrations, out of date extinguisher inspections.

Risk Factors

Tier	Driver
4	Recordable Incident Risk

Tripwires — the alert

Tripwires Risk Summary

Reporting Period: June 2023

SUFFOLK

7

NEWLY AT RISK

2

ESCALATING RISK

20

STABLE/IMPROVING

During July, 29 live wires were identified as at risk, including 7 newly at risk and 2 that escalated. The remaining 20 were stable or improving, and the overall risk position changed by 4 compared with the previous month.

NEWLY AT RISK (7)

ESCALATING RISK (2)

View Full Dashboard

Tripwires Dashboard

Reporting Period: July 2023

29 Projects at Risk

7

2

20

Showing data through July 2023 | Last Updated: August 10, 2023 | Download Report (PDF)

Project	Owner	Current Risk	Previous Risk	Change	Next Review	Notes
Project A	Owner A	High	High	Stable	2023-08-15	
Project B	Owner B	High	High	Stable	2023-08-15	
Project C	Owner C	High	High	Stable	2023-08-15	
Project D	Owner D	High	High	Stable	2023-08-15	
Project E	Owner E	High	High	Stable	2023-08-15	
Project F	Owner F	High	High	Stable	2023-08-15	
Project G	Owner G	High	High	Stable	2023-08-15	
Project H	Owner H	High	High	Stable	2023-08-15	
Project I	Owner I	High	High	Stable	2023-08-15	
Project J	Owner J	High	High	Stable	2023-08-15	
Project K	Owner K	High	High	Stable	2023-08-15	
Project L	Owner L	High	High	Stable	2023-08-15	
Project M	Owner M	High	High	Stable	2023-08-15	
Project N	Owner N	High	High	Stable	2023-08-15	
Project O	Owner O	High	High	Stable	2023-08-15	
Project P	Owner P	High	High	Stable	2023-08-15	
Project Q	Owner Q	High	High	Stable	2023-08-15	
Project R	Owner R	High	High	Stable	2023-08-15	
Project S	Owner S	High	High	Stable	2023-08-15	
Project T	Owner T	High	High	Stable	2023-08-15	
Project U	Owner U	High	High	Stable	2023-08-15	
Project V	Owner V	High	High	Stable	2023-08-15	
Project W	Owner W	High	High	Stable	2023-08-15	
Project X	Owner X	High	High	Stable	2023-08-15	
Project Y	Owner Y	High	High	Stable	2023-08-15	
Project Z	Owner Z	High	High	Stable	2023-08-15	


NEWLY AT RISK (7)

ESCALATING RISK (2)

View Full Dashboard


GETTING OFF THE PLATEAU

AQuA — the conversation



GETTING OFF THE PLATEAU

Vatic — judgment moves into the tool



GETTING OFF THE PLATEAU

Vera — the same idea, one step earlier





“Answers got cheap. The premium on the question went up.”

— *J.B.P. Shaw, “Is Your Cache Stale?”*

“The question” means both at once: the metric and the vehicle.

Don't let the medium become the mandate

Build for the problem. Let the shape follow.

Thank you

Questions?



@gottalottasoul



@abbenante

